

# Databases

## Part 1: Basics

Aleksandr Bernadskii

Solutions Architect  
Amazon Web Services



# There was a time when...

Mainframe



Client Server



Three tier



Microservices



# Dimensions to consider


Shape



Size



Compute



# Shape of data

Shape	Usecase	Example
Row store	Operate on a single or group of rows	Transactions, Payroll
Column store	Aggregations, big scans	Analytics, DSS
Key-Value	Low latency ingestion and query by Key	Tracking, Caching
Document	Index and store documents for query on attributes	Content Mgmt, Patient data
Graph	Store and query relationships	Recommendations
Time-series	Store and aggregate data sequenced entries	Telemetry, Sensors
Wide-column	Attribute-based data with query and sorting on columns	Features of a car
Blockchain	Ledger with immutable records	Audit records, Record history
Unstructured	Get and put Objects, Documents	Binary files, PDFs

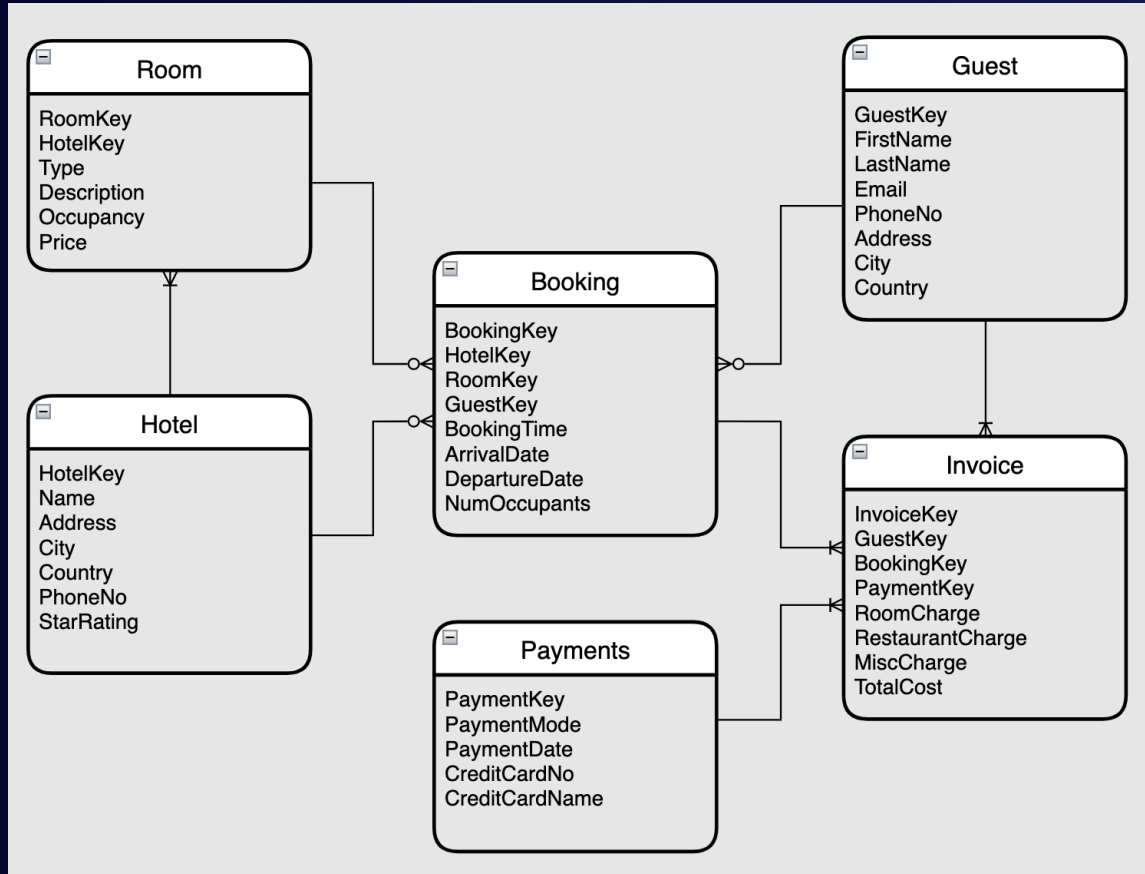
# Sizing and scaling

Considerations	Example
Size at limit – bound/unbound	Number of countries/cities – bound Number of sensors - unbounded
Working set size	Sales data history, but only 12 months relevant Session data of active users
Retrieval size and Caching	Get one row/document
Partitioned or monolithic	Sensor data is partitionable Company payroll does not have a natural partition

# Compute flexibility


Considerations	Example
Compute functions	Sum of sales for last 12 months Get and Put Transactions
Throughput	Milions of users getting hotels availability every minute Users scrolling a blog website
Latency	Get the location of a car every 5 seconds Get min,max of average bonus pay last year
Velocity of change	Inventory counts update frequently GPS data is almost never updated
Rate of ingestion	Device metrics data inserted every second Few new employees added every month

# Relational DBs or NoSQL?



```
{
  "object": {
    "type": "booking",
    "bookingKey": "asdbc-odpkg-otjg7",
    "bookingTime": "2022-05-09 10:00:00 CET",
    "arrivalDate": "2022-05-10 10:00:00 CET",
    "departureDate": "2022-05-12 10:00:00 CET",
    "rooms": [{
      "type": "DoubleRoom",
      "Description": " Double Room with see view",
      "price": 145,
      "hotel": {
        "name": "AnyHotel",
        "address": "AnyAddress",
        "city": "AnyCity",
        "country": "AnyCountry"
      }
    }]
  }
}
```

# CAP theorem. ACID and BASE



It is impossible to achieve both consistency and availability in a partition tolerant distributed system

ACID Model



BASE Model





# OLTP and OLAP

	OLTP	OLAP
Function	Day to day operation	Decision support
Database Design	Application oriented	Subject oriented
Data	Current, up-to-date detailed, flat relational, isolated	Historical, summarized multi-dimensional, consolidated
Usage	Repetitive	Ad-hoc
Access	Read/Write	Lots of scans
Unit of Work	Short, simple transaction	Complex query
Database Size	Gigabytes	Terabytes
Metric	Transaction throughput	Query throughput, response

# Relational Databases (RDB) vs Non-Relational Databases

RDB

NoSQL

Optimized for storage	Optimized for compute
Normalized/relational	Denormalized/hierarchical
Ad hoc queries	Instantiated views
Scale vertically	Scale horizontally
Good for OLAP, Limited Scale OLTP	Built for OLTP at scale

# SQL and NoSQL Scalability Pitfalls

- Scaling vertically is not infinitely
  - Adding more indexes increases the latency for insert/update
  - Transaction consistency overhead
  - Sharding may be difficult
  - Lock contention increases with the amount of transactions
- Not knowing your data access patterns
  - Adding too many indexes to satisfy new patterns
  - Sharding producing uneven or narrowed distribution of data
  - Not keeping related data together
  - Scanning whole items/tables instead of specific attributes

# Modern cloud-based applications

## Loosely coupled micro-services and purpose-built data stores

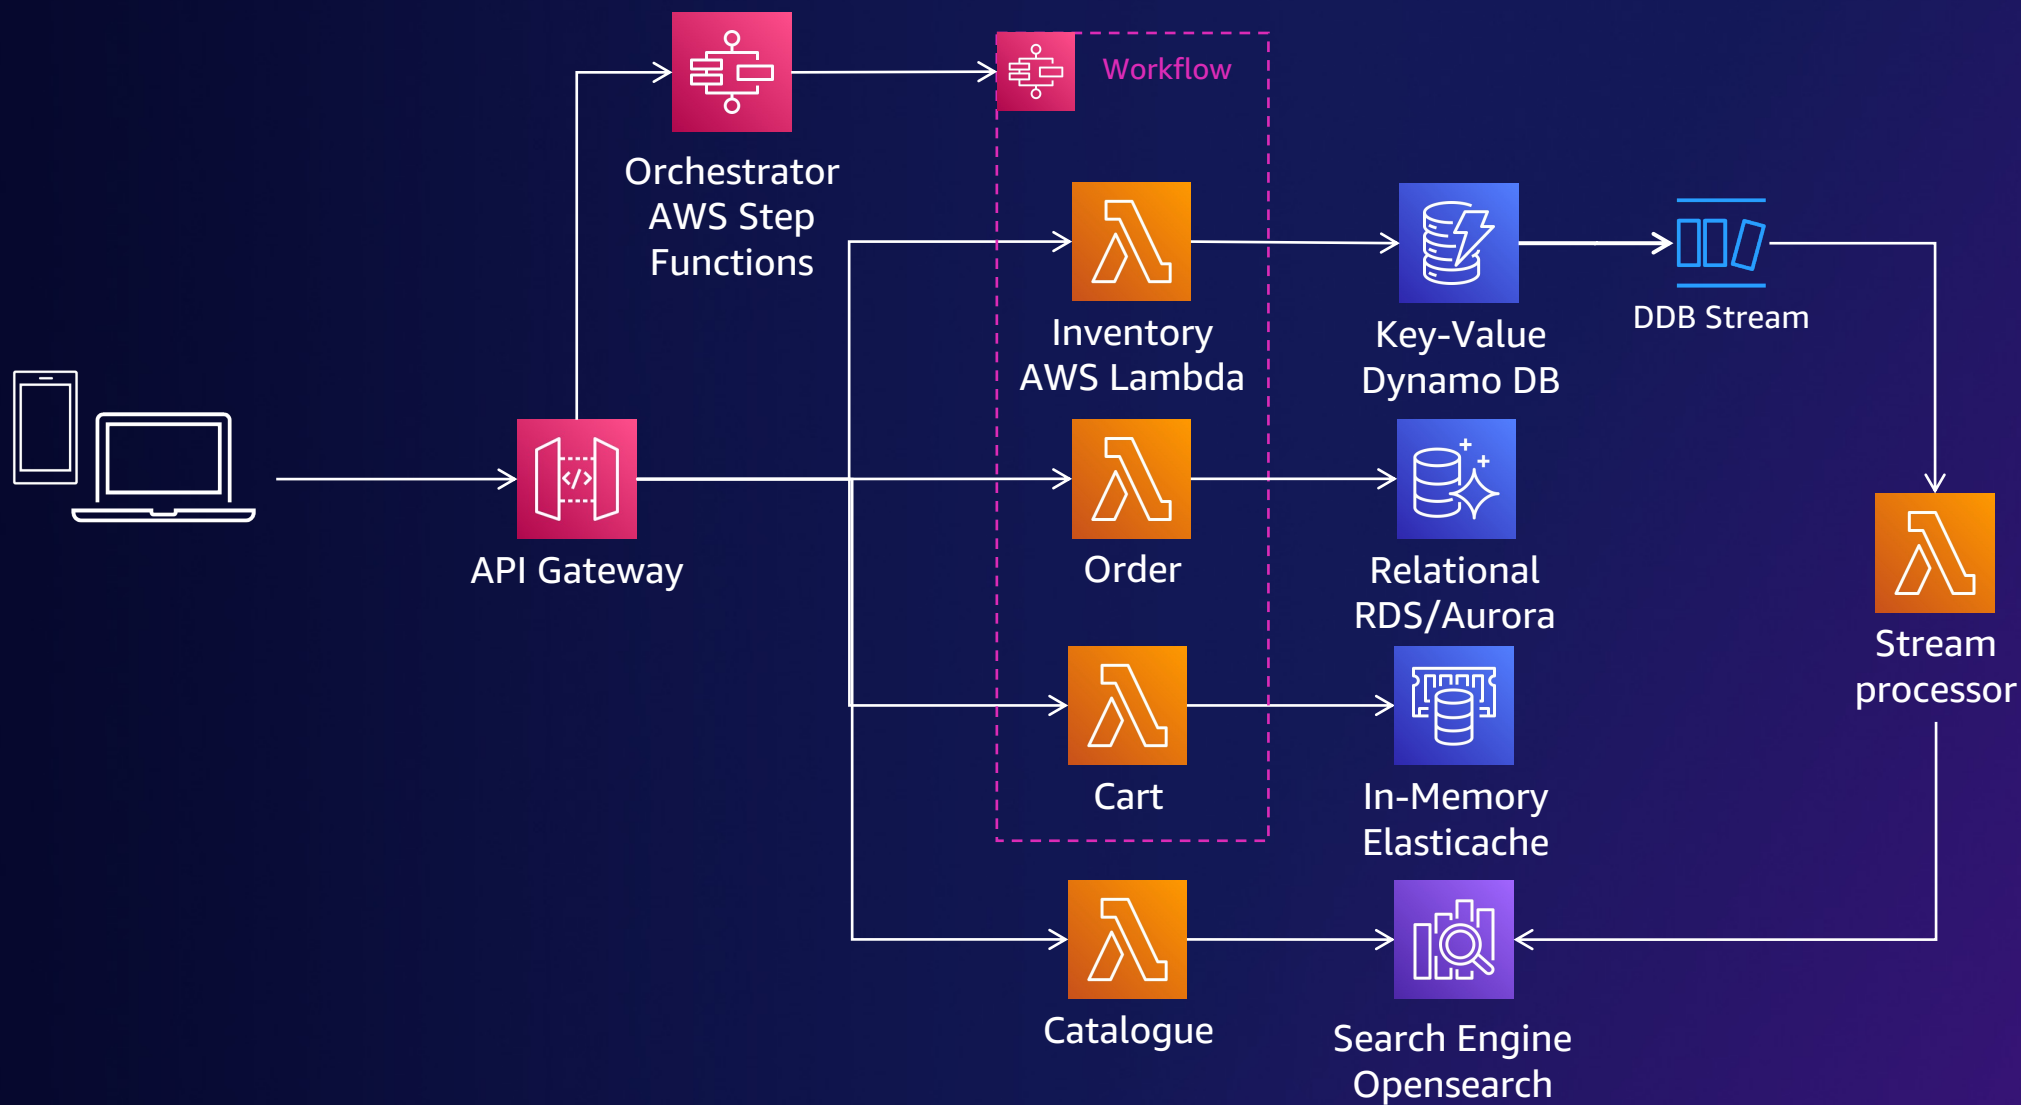


# AWSome pets - Current architecture



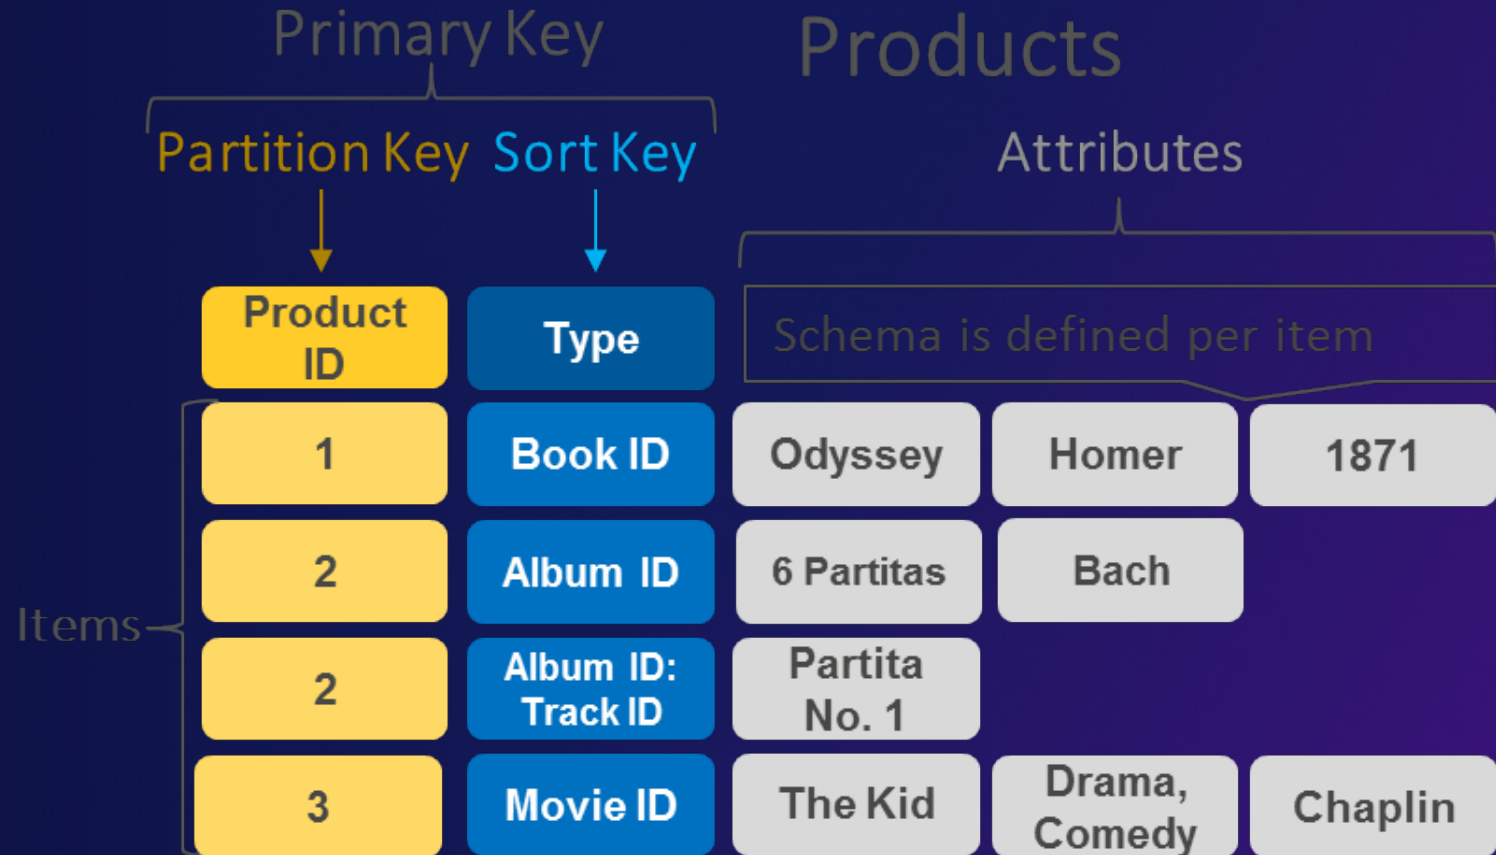
\* AWSome pets is a fictional company

# Target architecture

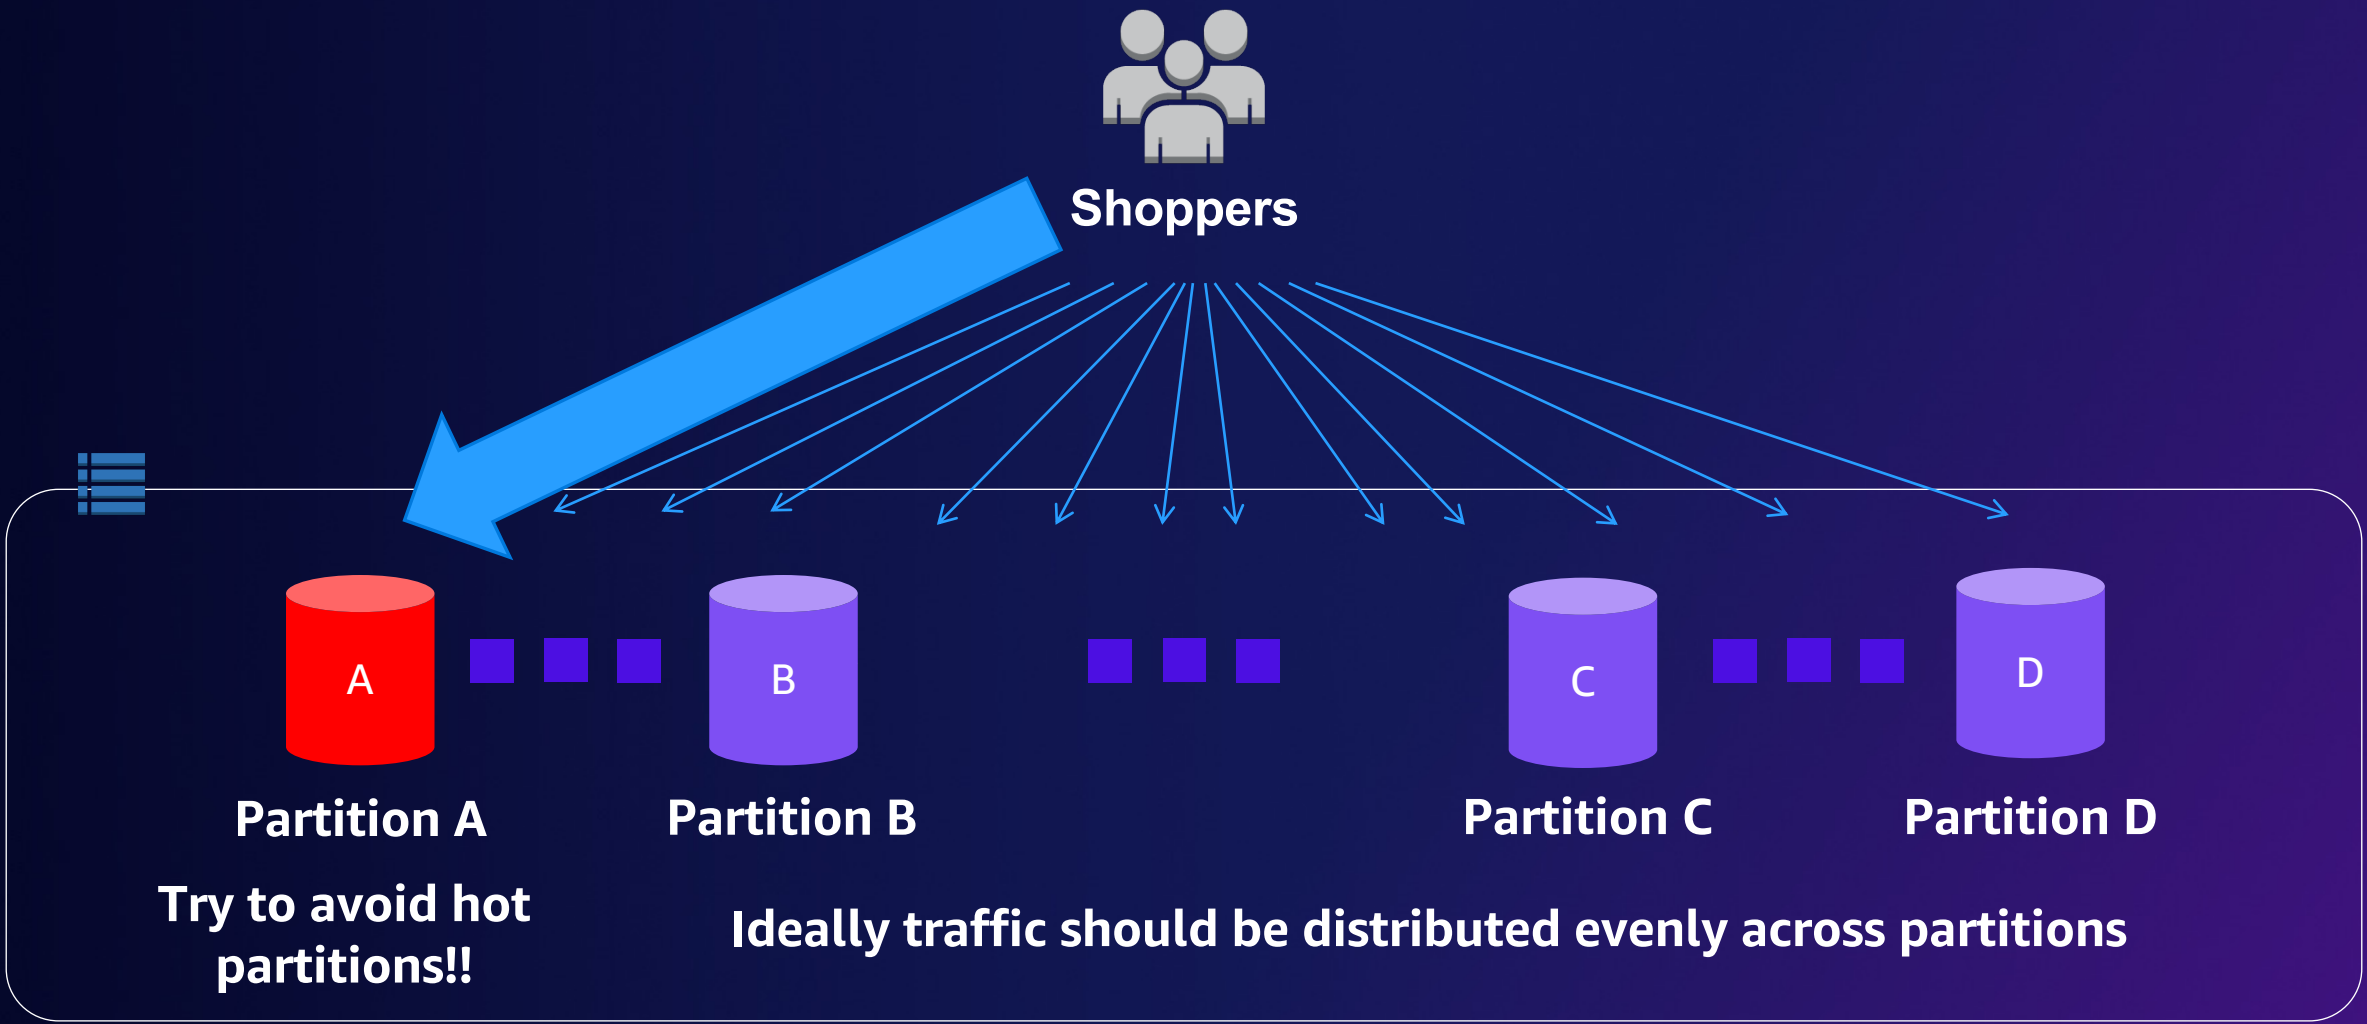


# Key-value database

- Simple key-value pairs
- Partitioned by keys
- Resilient to failure
- High-throughput, low-latency reads and writes
- Consistent performance at scale



# Good Partitioning





# Why document databases?

JSON is a flexible  
and natural  
format



JSON is the de  
facto API  
output



Store, query, &  
index JSON data  
natively



*JSON all the way*

# When should you use a document database?

JSON data



Flexible  
schema for  
fast iteration




Ad hoc query  
capabilities



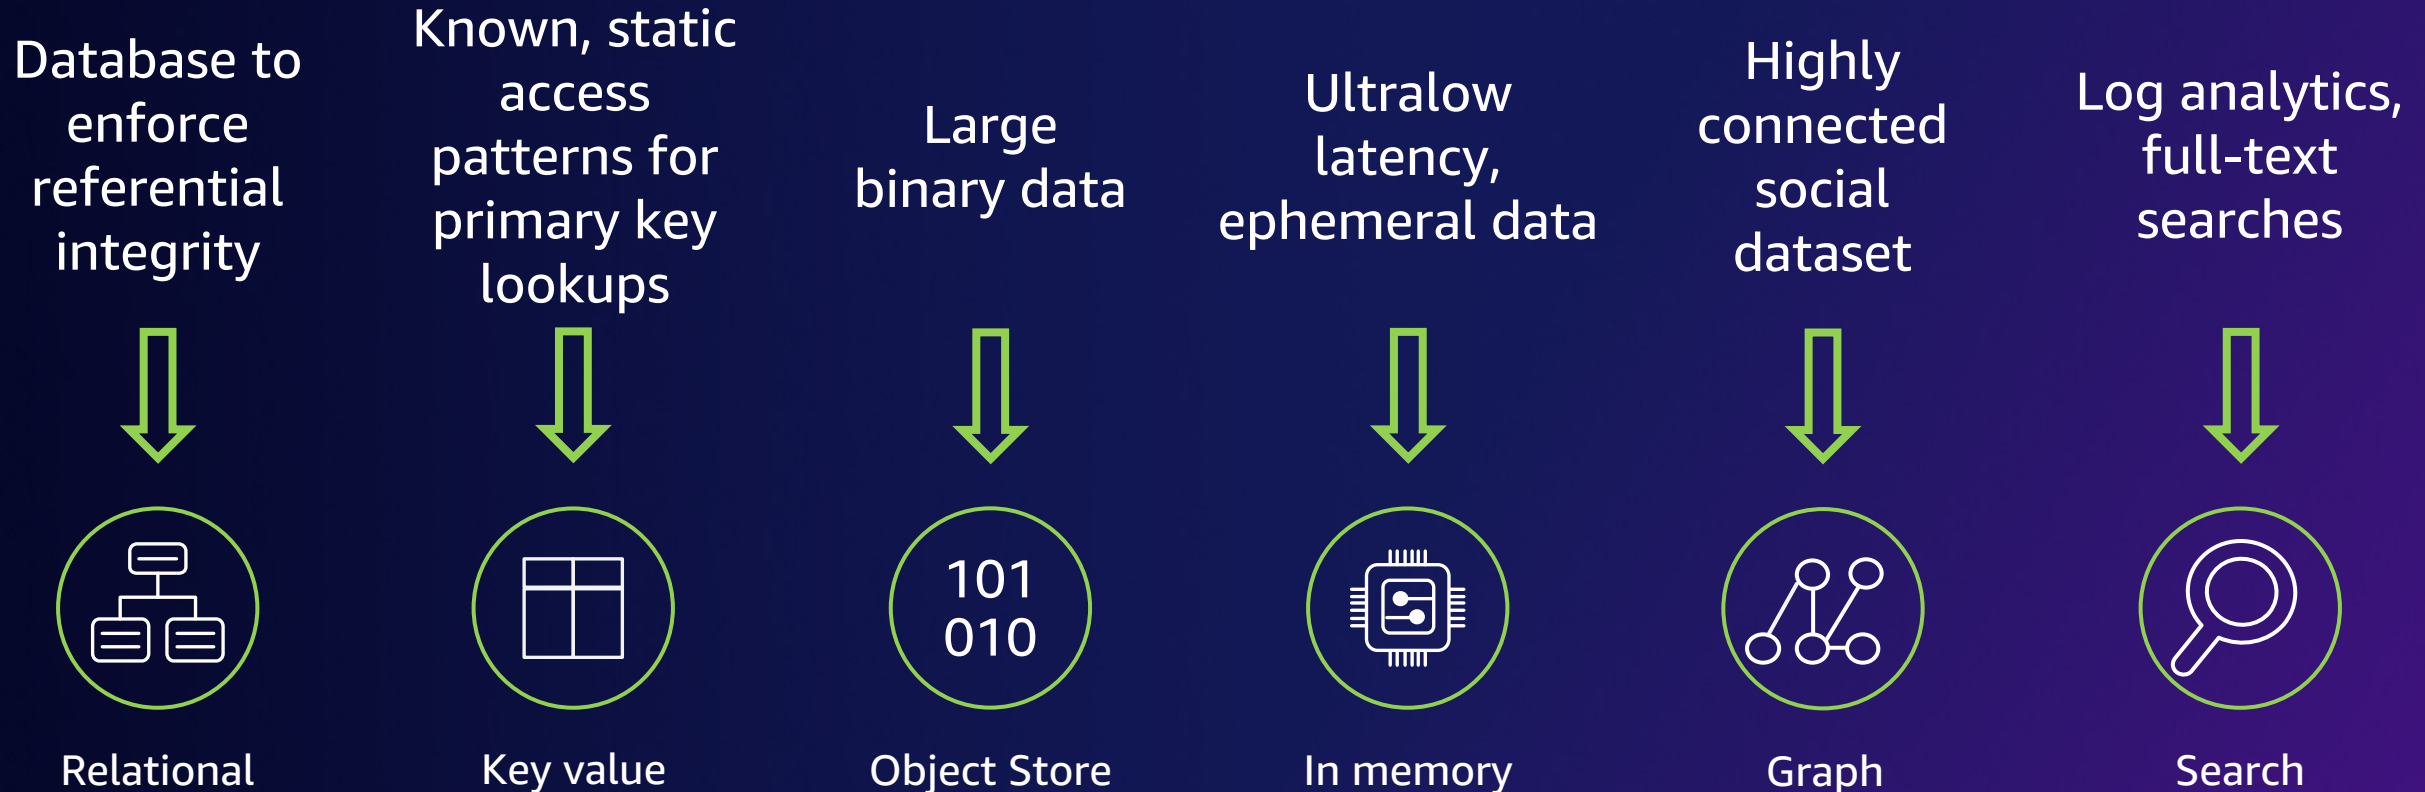
Flexible  
indexing



Operational and  
analytics  
workloads



# When are other databases more appropriate?



# Graph databases

Graph databases excel at answering questions like:


- What does this person want to buy?
- How are these two people connected?
- Why did X impact Y?

These questions:

- Navigate (variably) connected structure
- Filter or compute a result on the basis of the *strength*, *weight*, or *quality* of relationships
- Require traversing an unknown number of connections

# Query languages- Designed to move through data

- Graphs query languages are optimized to use connections to move through a network



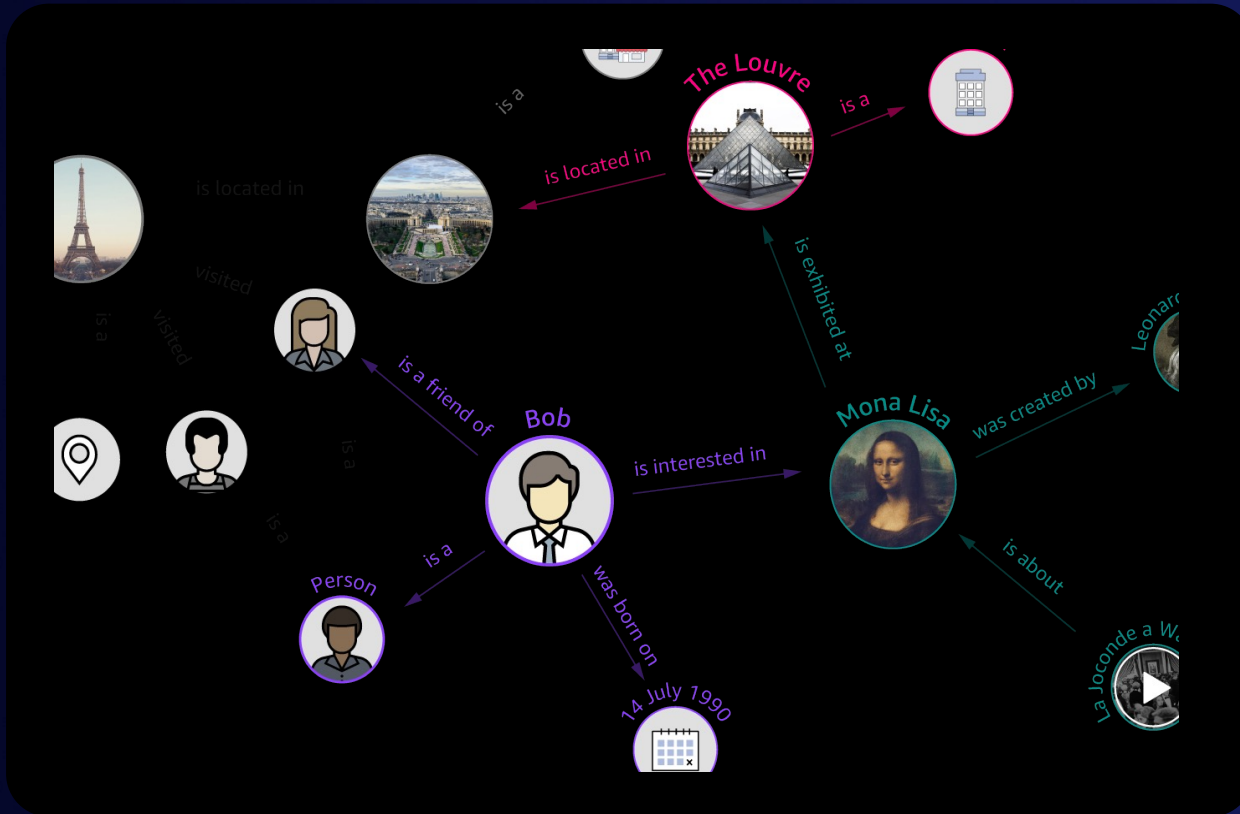
- Relational queries work by combining sets of data.

•

Person		Company			Person	Company
Andy	X	Amazon	=		Andy	Amazon

# What is a Knowledge Graph?

Understanding the who, what, when, and where



## Benefits

### 1. Link disparate data sources

Link disparate and heterogeneous data sources together to discover hidden connections

### 2. Improved search results

Increase productivity by making data easily accessible through improved search relevance

### 3. Augment ML/AI

Improve the efficiency and effectiveness of machine learning models by providing context and augmentation with related content

<https://aws.amazon.com/blogs/apn/exploring-knowledge-graphs-on-amazon-neptune-using-metaphactory/>

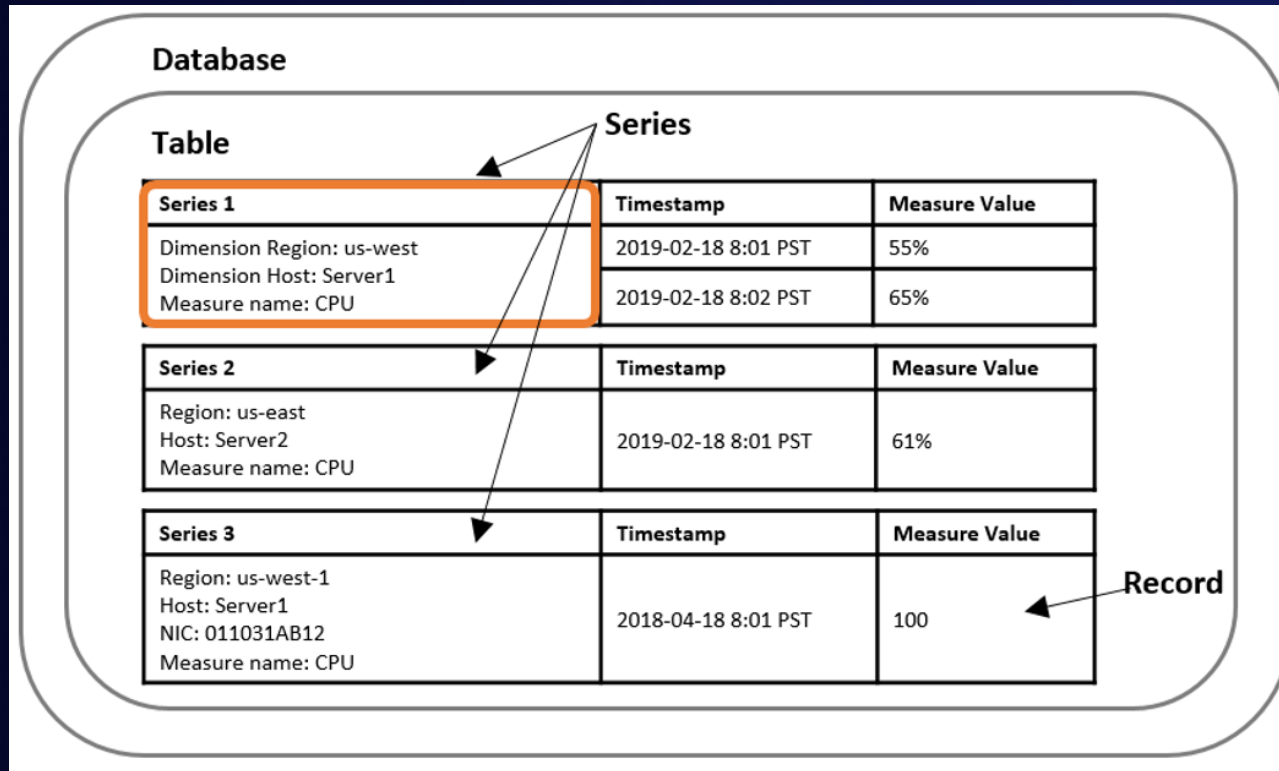
# Characteristics of time-series data

A sequence of data points recorded over a time interval

- Data is append-only
- Data typically arrives in time-order
- Queries are over a time interval
- Specialized math is required to find trends and patterns
- Older data is typically aggregated over time or discarded




















# Timeseries databases. Concepts



- **Series** - A sequence of data points recorded over a time interval
- **Record** - A single data point in the time series.
- **Dimension** - An attribute that describes the meta-data of the time series.
- **Measure** - An attribute that describes the data of the series
- **Timestamp** - Every record has a timestamp, which indicates when the **Measure** was collected.
- **Table** - A table is a container for (related) time series with timestamp, dimensions and measures.
- **Retention policy** - A value you can set at the table level to dictate how long time series data will be retained.
- **Database** - A top level container for tables and retention policies.



# AWS Purpose built databases

								
<b>Relational</b>	<b>Key-value</b>	<b>Document</b>	<b>In-memory</b>	<b>Graph</b>	<b>Time-series</b>	<b>Ledger</b>	<b>Wide column</b>	
Referential integrity, ACID transactions, schema-on-write	High throughput, low-latency reads and writes, endless scale	Store documents; quickly access querying on any attribute	Query by key with microsecond latency	Quickly and easily create and navigate relationships between data	Collect, store, and process data sequenced by time	Complete, immutable and verifiable history of all changes to application data	Scalable, highly available, and managed Apache Cassandra-compatible service	
 Aurora	 RDS	 DynamoDB	 DocumentDB	 ElastiCache	 Neptune	 Timestream	 QLDB	 Amazon Keyspaces
Lift and shift, ERP, CRM, finance	Real-time bidding, shopping cart, social, product catalog, customer preferences	Content management, personalization, mobile	Leaderboards, real-time analytics, caching	Fraud detection, social networking, recommendation engine	IoT applications, event tracking	Systems of record, supply chain, health care, registrations, financial	Build low-latency applications, leverage open source, migrate Cassandra to the cloud	

# Thank you!



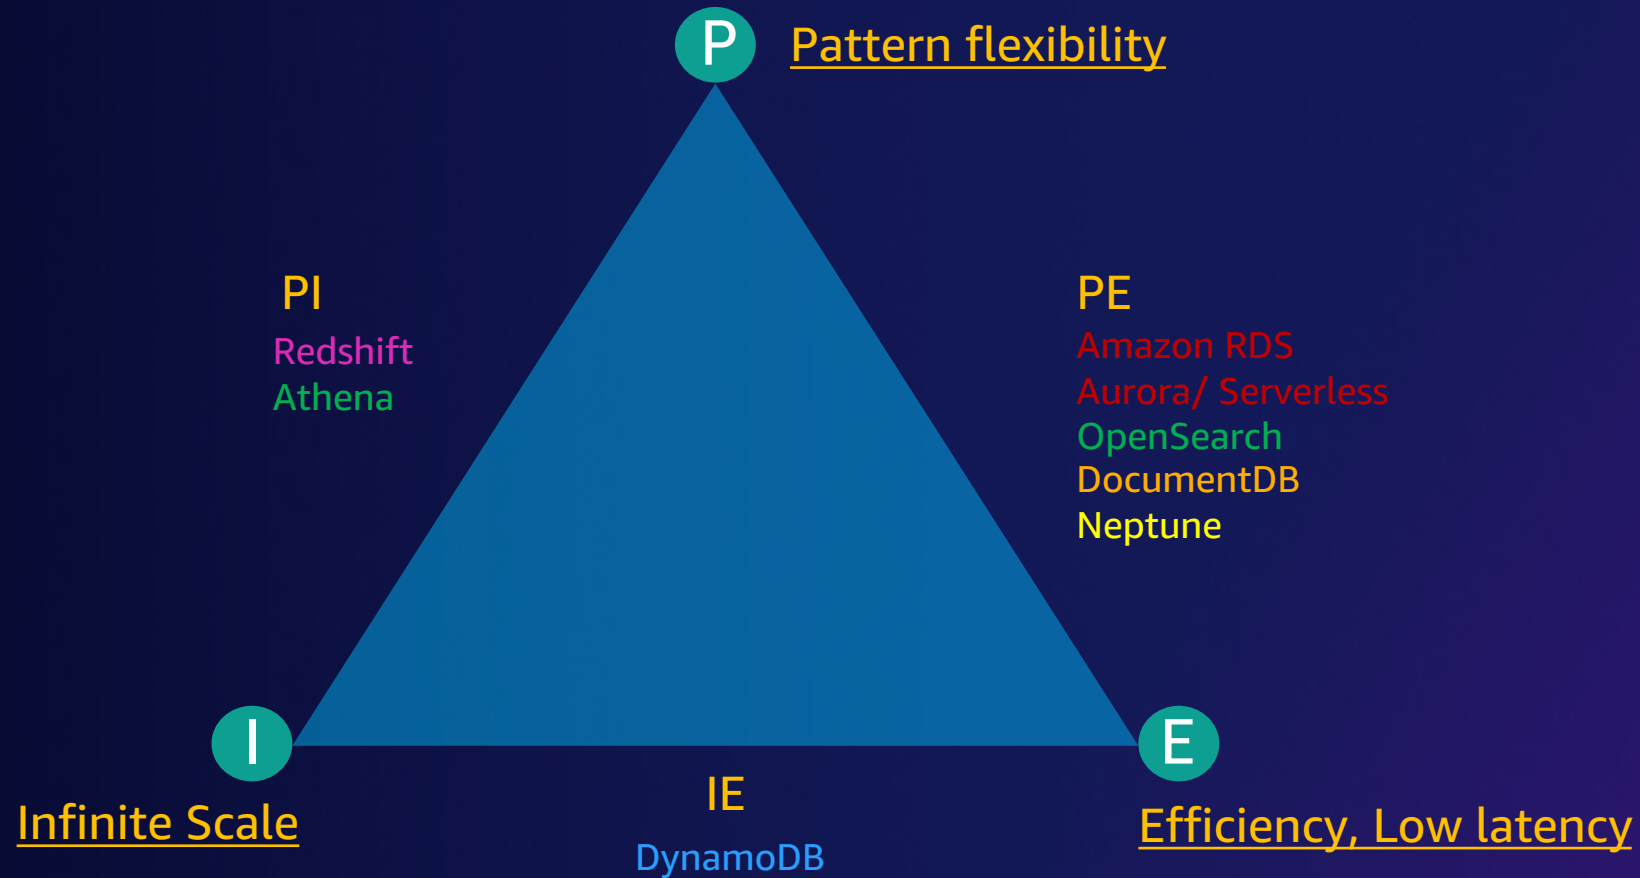
# Decision Matrix

Service	Shape	Size	Workload	Performance	Durability	Expertise	Legacy	Business Needs	Ops. Load	Platform Integration
Amazon Aurora MySQL	Structured, semistruct.	Mid TB Range	K/V lookups, transactional, light analytics	High throughput, low latency	High	Relational, MySQL, SQL Server	User defined code, COTS	Agility, cost opt.	Low to moderate	Serverless, IAM, Lambda, Auto Scaling, Amazon S3
Amazon Aurora PostgreSQL	Structured, semistruct.	Mid TB Range	Transactional, light analytics	High throughput, low latency	High	Relational, PostgreSQL, Oracle	User defined code, COTS	Agility, cost opt.	Moderate	In the works
Amazon RDS DB Engines	Structured, semistruct.	Low TB Range	Transactional, light analytics	Mid-to-high throughput, low latency	User controlled	Engine specific	Engine req., COTS	Right sizing	Moderate	Log streaming
Amazon DynamoDB	Semistruct.	High TB Range	K/V lookups, NoSQL, OLTP, document store	Ultra-high throughput, low latency	High	NoSQL	No active development	Zero downtime, Ultra-high scale	Low	Serverless, IAM, Lambda, DAX, Auto Scaling, Kinesis Streams
Amazon Neptune	Graph-structured	Mid TB Range	Graph, highly connected data, transact.	High throughput, low latency	High	Graph, Gremlin, SPARQL		Agility, cost opt.	Low	IAM, Amazon S3

# Decision Matrix (cont.)

Service	Shape	Size	Workload	Performance	Durability	Expertise	Legacy	Business Needs	Ops. Load	Platform Integration
<b>Amazon ElastiCache</b>	Semistruct., Unstructured	Low TB Range	In-memory caching, K/V lookups, NoSQL	High throughput, ultra-low latency	Low (In-memory, auto failover to read replica)	Caching, NoSQL		Response latency, DB cost opt.	Low	Scalable clusters
<b>Amazon Redshift</b>	Structured, semistruct.	PB Range	Optimized analytics	Mid-to-high latency	High	DW, data science	User defined code, COTS	Cost opt.	Moderate	IAM, Amazon S3
<b>Amazon Athena</b>	Structured, semistruct.	TB Range	Flexible analytics	High latency	High (Amazon S3)	Data lakes, data science		Flexibility	Low	IAM, Amazon S3
<b>Amazon EMR</b>	Semistruct.	PB Range	Flexible analytics	Low-to-high latency	User-controlled	Data lakes, data science, DW	Tooling & versioning	Cost opt.	High	IAM, Amazon S3, EC2 Spot

# PIE Theorem for data stores



## Data Models:

Relational

Columnar

Wide column

Document

Graph

Unstructured